

PREFIX SUM REPORT

Name: Udayin Biswas

Entry Number: 2014CS50296

Machine Specifications: 2.5 GHz I5 core(quad core), 4 gb ram

For calculating the Prefix Sum, it had to be done in $O(\log n)$ phases at $O(n)$ cost.

This calculation was done in a specific manner as follows:

We are provided with an input vector of the numbers and suppose n threads.

Assumption: n is 2 to the power of k , $k > 0$.

Data structure used was a vector(/array) which we can assume to take the shape of a tree when computation is performed in $\log n$ phases.

First, we have to sum up certain elements at each stage, with $\log(n)$ (to the base 2) stages.

At stage 0, every 2 neighbour elements at a gap of 1 get summed up together .

$a[1]=a[0]+a[1]$

$a[3]=a[2]+a[3]$

$a[5]=a[4]+a[5]$

$a[7]=a[7]+a[6]$

At stage 1, every 2 elements at a gap of 2 get summed up.

$a[3]=a[3]+a[1]$

$a[7]=a[7]+a[5]$

At stage 2, every 4 elements at a gap of 4 get summed up.

$a[7]=a[7]+a[3]$

....

and so on.

So total we can see we have $\log(n)$ phases, and the computation cost halves at each increasing stage.

Eg: 1,2,3,4,5,6,7,8

->Stage 0: 1,3,3,7,5,11,7,15

->Stage 1: 1,3,3,10,5,11,7,26

->Stage 2: 1,3,3,10,5,11,7,36

Now next stage,

we have to traverse back from $a[7]$ to rest of the elements in a manner.

first: the last element is set to 0.

-> 1,3,3,10,5,11,7,36

At each stage now, left child value = temp, left child value=right child value, right child value=temp +left child value. $\log n$ such phases again.

Eg above- 1,3,3,10,5,11,7,0

Stage 0: $a[3]=0, a[7]=9+0=9$ -> 1,3,3,0,5,11,7,9

Stage 1: $a[1]=0, a[3]=3, a[5]=9, a[7]=14$ -> 1,0,3,6,5,9,7,14

and so on,


we finally get 0,1,3,6,10,15,21,28

Shift by 1 to the left and calculate the final value by adding 8 to 28.

=>1,3,6,10,15,21,28,36. So the basic strategy was to first sum up elements region wise so that we can get the total at the end. then we set it 0 so that 0 can travel to the left. Meanwhile the right child is the sum of itself and the left child value at each stage. The left child carries the 0 to the left most of the vector. This strategy will help us get prefix sum.

Parallelisation was done on the 2 for loops to divide elements of array among the threads, at each stage. Each thread takes the next consecutive element at a gap of n threads for computation so that the data is divided equally among the threads.

Speedup vs problem size



Efficiency vs problem size

